

Assignment 3: Feature Stores

Credits: I used the Amazon Documentation Developer Guide to help with the feature store. Link provided [HERE](#)

Creating Feature Store Session

```
In [ ]: # Import the necessary libraries
# Libraries for creating feature store session
import boto3
import sagemaker
from sagemaker.session import Session
from sagemaker import get_execution_role

# Libraries for interacting with the dataset
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import io

# Feature Group
import time
from time import gmtime, strftime, sleep
from sagemaker.feature_store.feature_group import FeatureGroup
```

```
sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml
sagemaker.config INFO - Not applying SDK defaults from location: /root/.config/sagemaker/config.yaml
```

```
In [ ]: # Helpful Functions + Variables stored here
def encode_col(df, col):
    names = df[col].unique()
    values = len(names)
    dict_pairs = dict([(key,value) for _, (key,value) in enumerate(zip(names,np.arange(values)))]])
    df[col] = df[col].map(dict_pairs)
    df[col] = df[col].astype('float64')
    return df, dict_pairs

# def encode_col(df, col):
#     df[col], _ = df[col].factorize()
```

```
# df[col] = df[col].astype('float64')
# return df

def convert_to_strings(df):
    for col in df.columns:
        if df.dtypes[col] == 'object':
            df[col] = df[col].astype(str)

def wait_for_feature_group_creation_complete(feature_group):
    status = feature_group.describe().get("FeatureGroupStatus")
    while status == "Creating":
        print("Waiting for Feature Group Creation")
        time.sleep(5)
        status = feature_group.describe().get("FeatureGroupStatus")
    if status != "Created":
        raise RuntimeError(f"Failed to create feature group {feature_group.name}")
    print(f"FeatureGroup {feature_group.name} successfully created.")

time_now = int(round(time.time()))
```

```
In [ ]: # Create the session by identifying the variables
region = boto3.Session().region_name

boto_session = boto3.Session(region_name=region)

sagemaker_client = boto_session.client(service_name="sagemaker", region_name=region)
featurestore_runtime = boto_session.client(
    service_name="sagemaker-featurestore-runtime", region_name=region
)

feature_store_session = Session(
    boto_session=boto_session,
    sagemaker_client=sagemaker_client,
    sagemaker_featurestore_runtime_client=featurestore_runtime,
)
```

```
In [ ]: # Creating default bucket
default_s3_bucket_name = feature_store_session.default_bucket()
prefix = "sagemaker-featurestore-demo"
```

```
print(default_s3_bucket_name)
```

sagemaker-us-east-1-004608622582

```
In [ ]: # Grab Role
        role = get_execution_role()
        print(role)
```

arn:aws:iam::004608622582:role/LabRole

```
In [ ]: # Start the client + feature store runtime
        sagemaker_client = boto_session.client(service_name='sagemaker', region_name=region)
        featurestore_runtime = boto_session.client(service_name='sagemaker-featurestore-runtime', region_name=region)
```

```
In [ ]: # Create feature store session
        feature_store_session = Session(boto_session=boto_session, sagemaker_client=sagemaker_client, sagemaker_featurestore_
```

Loading Data and Partitioning it into DataGroups

```
In [ ]: # Reading in the data
        housing_df = pd.read_csv('housing.csv')
        housing_gmaps_df = pd.read_csv('housing_gmaps_data_raw.csv')
```

```
In [ ]: housing_df.head()
```

```
Out[ ]: 
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_hou
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	



```
In [ ]: housing_gmaps_df.head()
```

Out[]:

	street_number	route	locality-political	administrative_area_level_2-political	administrative_area_level_1-political	country-political	postal_code	address
0	3130	Grizzly Peak Boulevard	Berkeley	Alameda County	California	United States	94705.0	3130 Grizzly Peak Blvd, Berkeley, CA 94705, USA
1	2005	Tunnel Road	Oakland	Alameda County	California	United States	94611.0	2005 Tunnel Rd, Oakland, CA 94611, USA
2	6886	Chabot Road	Oakland	Alameda County	California	United States	94618.0	6886 Chabot Rd, Oakland, CA 94618, USA
3	6365	Florio Street	Oakland	Alameda County	California	United States	94618.0	6365 Florio St, Oakland, CA 94618, USA
4	5407	Bryant Avenue	Oakland	Alameda County	California	United States	94618.0	5407 Bryant Ave, Oakland,

street_number	route	locality-political	administrative_area_level_2-political	administrative_area_level_1-political	country-political	postal_code	address
							CA 94618, USA

5 rows × 30 columns

In []: `housing_df.describe()`

Out []:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_inc
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870000
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.890000
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.490000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.560000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.530000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.740000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000000



In []: `housing_gmaps_df.describe()`

Out[]:

	postal_code	longitude	latitude	postal_code_suffix
count	12410.000000	12590.000000	12590.000000	7999.000000
mean	93348.943836	-119.676724	35.895577	4177.914614
std	1765.572652	2.042677	2.219248	2474.063791
min	85344.000000	-124.350000	32.540000	110.000000
25%	92054.000000	-121.760000	33.970000	2230.500000
50%	93301.000000	-119.270000	35.340000	3556.000000
75%	95050.000000	-117.950000	37.810000	5529.000000
max	96161.000000	-114.310000	41.950000	9859.000000

```
In [ ]: df = pd.merge(housing_gmaps_df, housing_df, left_on=['longitude', 'latitude'], right_on=['longitude', 'latitude'], ho
```

```
In [ ]: df.head()
```

Out[]:

	street_number	route	locality-political	administrative_area_level_2-political	administrative_area_level_1-political	country-political	postal_code	address
0	3130	Grizzly Peak Boulevard	Berkeley	Alameda County	California	United States	94705.0	3130 Grizzly Peak Blvd, Berkeley, CA 94705, USA
1	2005	Tunnel Road	Oakland	Alameda County	California	United States	94611.0	2005 Tunnel Rd, Oakland, CA 94611, USA
2	6886	Chabot Road	Oakland	Alameda County	California	United States	94618.0	6886 Chabot Rd, Oakland, CA 94618, USA
3	6365	Florio Street	Oakland	Alameda County	California	United States	94618.0	6365 Florio St, Oakland, CA 94618, USA
4	6365	Florio Street	Oakland	Alameda County	California	United States	94618.0	6365 Florio St, Oakland, CA

street_number	route	locality-political	administrative_area_level_2-political	administrative_area_level_1-political	country-political	postal_code	address
							94618, USA

5 rows × 38 columns

```
In [ ]: # Priority Key
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 20640 entries, 0 to 20639
```

```
Data columns (total 38 columns):
```

#	Column	Non-Null Count	Dtype
0	street_number	19008 non-null	object
1	route	20091 non-null	object
2	locality-political	20452 non-null	object
3	administrative_area_level_2-political	20589 non-null	object
4	administrative_area_level_1-political	20637 non-null	object
5	country-political	20640 non-null	object
6	postal_code	20454 non-null	float64
7	address	20640 non-null	object
8	longitude	20640 non-null	float64
9	latitude	20640 non-null	float64
10	neighborhood-political	9000 non-null	object
11	postal_code_suffix	14095 non-null	float64
12	establishment-point_of_interest-transit_station	255 non-null	object
13	establishment-park-point_of_interest	46 non-null	object
14	premise	36 non-null	object
15	establishment-point_of_interest-subway_station-transit_station	3 non-null	object
16	airport-establishment-finance-moving_company-point_of_interest-storage	1 non-null	object
17	subpremise	25 non-null	object
18	bus_station-establishment-point_of_interest-transit_station	22 non-null	object
19	establishment-park-point_of_interest-tourist_attraction	34 non-null	object
20	establishment-natural_feature	11 non-null	object
21	airport-establishment-point_of_interest	8 non-null	object
22	political-sublocality-sublocality_level_1	33 non-null	object
23	administrative_area_level_3-political	1 non-null	object
24	post_box	6 non-null	object
25	establishment-light_rail_station-point_of_interest-transit_station	13 non-null	object
26	establishment-point_of_interest	1 non-null	object
27	aquarium-establishment-park-point_of_interest-tourist_attraction-zoo	1 non-null	object
28	campground-establishment-lodging-park-point_of_interest-rv_park-tourist_attraction	1 non-null	object
29	cemetery-establishment-park-point_of_interest	1 non-null	object
30	housing_median_age	20640 non-null	float64
31	total_rooms	20640 non-null	float64
32	total_bedrooms	20433 non-null	float64
33	population	20640 non-null	float64
34	households	20640 non-null	float64
35	median_income	20640 non-null	float64
36	median_house_value	20640 non-null	float64

```
37 ocean_proximity
dtypes: float64(11), object(27)
memory usage: 6.0+ MB
```

20640 non-null object

```
In [ ]: # Grabbing the features for our group
feature_cols = ['neighborhood-political',
                'ocean_proximity',
                'median_house_value',
                'housing_median_age',
                'households',
                'total_bedrooms',
                'locality-political']

# Creating new df based on the destired features
feature_df = df[feature_cols]

# Dropping an null values based on the primary key
feature_df = feature_df.dropna(subset='neighborhood-political')

# Renaming some of the columss for simplicity sake
feature_df = feature_df.rename(columns={'neighborhood-political': 'nbh_pol',
                                       'locality-political': 'loc_pol',
                                       'ocean_proximity': 'ocn_prox',
                                       'median_house_value': 'med_hse_val',
                                       'housing_median_age': 'hse_med_age',
                                       'households': 'tot_house',
                                       'total_bedrooms': 'tot_bed'})
```

```
In [ ]: feature_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 9000 entries, 1 to 20636
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   nbh_pol      9000 non-null   object
1   ocn_prox     9000 non-null   object
2   med_hse_val  9000 non-null   float64
3   hse_med_age  9000 non-null   float64
4   tot_house    9000 non-null   float64
5   tot_bed      8911 non-null   float64
6   loc_pol      8955 non-null   object
dtypes: float64(4), object(3)
memory usage: 562.5+ KB
```

```
In [ ]: # Households because on Locality
house_df = feature_df[['loc_pol',
                       'tot_house',
                       'tot_bed']]

# Finding average for Locality Code
house_df = house_df.groupby('loc_pol').mean()

# Renaming total to averages for average colculation
house_df = house_df.rename(columns={'tot_bed': 'avg_bed',
                                   'tot_house': 'avg_house'})

# Finding average bedrooms per household
house_df['avg_bed_per_house'] = house_df['avg_bed'].div(house_df['avg_house'], axis=0).round()

# Merge the new df back into the feature_df
feature_df = pd.merge(feature_df, house_df, left_on=['loc_pol'], right_on=['loc_pol'], how='left')

In [ ]: feature_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9000 entries, 0 to 8999
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   nbh_pol                9000 non-null   object
1   ocn_prox               9000 non-null   object
2   med_hse_val            9000 non-null   float64
3   hse_med_age            9000 non-null   float64
4   tot_house              9000 non-null   float64
5   tot_bed                8911 non-null   float64
6   loc_pol                8955 non-null   object
7   avg_house              8955 non-null   float64
8   avg_bed                8954 non-null   float64
9   avg_bed_per_house      8954 non-null   float64
dtypes: float64(7), object(3)
memory usage: 703.2+ KB
```

```
In [ ]: # Encoding datatypes from objects to floats

# One hot encode ocn_prox
encode_ocn_prox = pd.get_dummies(feature_df['ocn_prox'], dtype='float64')

# Encode Locality
encode_loc_pol, dict_pairs = encode_col(feature_df, 'loc_pol')

# Combine the two encoded columns together
combined_cols = pd.concat([encode_loc_pol, encode_ocn_prox, ], axis=1)
```

```
In [ ]: combined_cols.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9000 entries, 0 to 8999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   nbh_pol                9000 non-null   object
1   ocn_prox               9000 non-null   object
2   med_hse_val            9000 non-null   float64
3   hse_med_age            9000 non-null   float64
4   tot_house              9000 non-null   float64
5   tot_bed                8911 non-null   float64
6   loc_pol                9000 non-null   float64
7   avg_house              8955 non-null   float64
8   avg_bed                8954 non-null   float64
9   avg_bed_per_house      8954 non-null   float64
10  <1H OCEAN              9000 non-null   float64
11  INLAND                 9000 non-null   float64
12  NEAR BAY               9000 non-null   float64
13  NEAR OCEAN             9000 non-null   float64
dtypes: float64(12), object(2)
memory usage: 984.5+ KB
```

```
In [ ]: # Now that we got what we need form ocn_prox, we can drop that column
combined_cols = combined_cols.drop(columns=['ocn_prox'])
```

```
In [ ]: combined_cols.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9000 entries, 0 to 8999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   nbh_pol                9000 non-null   object
1   med_hse_val            9000 non-null   float64
2   hse_med_age            9000 non-null   float64
3   tot_house              9000 non-null   float64
4   tot_bed                8911 non-null   float64
5   loc_pol                9000 non-null   float64
6   avg_house              8955 non-null   float64
7   avg_bed                8954 non-null   float64
8   avg_bed_per_house      8954 non-null   float64
9   <1H OCEAN              9000 non-null   float64
10  INLAND                 9000 non-null   float64
11  NEAR BAY               9000 non-null   float64
12  NEAR OCEAN             9000 non-null   float64
dtypes: float64(12), object(1)
memory usage: 914.2+ KB
```

```
In [ ]: # Rename and Group the neighborhoods and create an index out of them
new_df = combined_cols.groupby('nbh_pol').mean()
```

```
In [ ]: new_df.head()
```

Out[]:

	med_hse_val	hse_med_age	tot_house	tot_bed	loc_pol	avg_house	avg_bed	avg_bed_per_house	<1H OCEAN	II
nbh_pol										
28 Palms	222200.000000	25.0	923.000000	939.000000	5.0	863.238806	894.686567	1.0	1.0	
Acorn Industrial	81300.000000	52.0	147.000000	244.000000	0.0	370.966197	397.541076	1.0	0.0	
Adams Hill	250733.333333	39.5	493.666667	520.166667	36.0	579.542056	614.600000	1.0	1.0	
Agua Mansa Industrial Corridor	112300.000000	17.0	516.000000	569.000000	138.0	516.000000	569.000000	1.0	0.0	
Al Tahoe	109180.000000	23.8	248.800000	399.800000	20.0	248.800000	399.800000	2.0	0.0	

In []:

```
# Rename
new_df = new_df.reset_index().rename(columns={'index': 'nbh_pol',
                                             '<1H OCEAN': 'ls_1_ocn',
                                             'INLAND': 'inland',
                                             'NEAR BAY': 'nr_bay',
                                             'NEAR OCEAN': 'nr_ocn'})
```

In []:

```
new_df.head()
```

Out[]:

	nbh_pol	med_hse_val	hse_med_age	tot_house	tot_bed	loc_pol	avg_house	avg_bed	avg_bed_per_house	ls_1_ocn
0	28 Palms	222200.000000	25.0	923.000000	939.000000	5.0	863.238806	894.686567	1.0	1.0
1	Acorn Industrial	81300.000000	52.0	147.000000	244.000000	0.0	370.966197	397.541076	1.0	0.0
2	Adams Hill	250733.333333	39.5	493.666667	520.166667	36.0	579.542056	614.600000	1.0	1.0
3	Agua Mansa Industrial Corridor	112300.000000	17.0	516.000000	569.000000	138.0	516.000000	569.000000	1.0	0.0
4	Al Tahoe	109180.000000	23.8	248.800000	399.800000	20.0	248.800000	399.800000	2.0	0.0



In []: new_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1306 entries, 0 to 1305
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   nbh_pol                1306 non-null  object
1   med_hse_val            1306 non-null  float64
2   hse_med_age            1306 non-null  float64
3   tot_house              1306 non-null  float64
4   tot_bed                1300 non-null  float64
5   loc_pol                1306 non-null  float64
6   avg_house              1293 non-null  float64
7   avg_bed                1292 non-null  float64
8   avg_bed_per_house      1292 non-null  float64
9   ls_1_ocn              1306 non-null  float64
10  inland                 1306 non-null  float64
11  nr_bay                 1306 non-null  float64
12  nr_ocn                 1306 non-null  float64
dtypes: float64(12), object(1)
memory usage: 132.8+ KB
```



```
In [ ]: # Encode the nbh_pol
# Creating a dataframe
encode_df = pd.DataFrame.from_dict(dict_pairs, orient='index')

# Reset the index and name the column
encode_df = encode_df.reset_index().rename(columns={'index': 'nbh_pol_new',
                                                    0: 'nbh_pol_encode'})

# encode_df.info()

# Ensure DF for the encoded values are the same
encode_df['nbh_pol_encode'] = encode_df['nbh_pol_encode'].astype('float64')
```

```
In [ ]: encode_df
```

```
Out[ ]:
```

	nbh_pol_new	nbh_pol_encode
0	Oakland	0.0
1	Berkeley	1.0
2	San Leandro	2.0
3	Alameda	3.0
4	Hayward	4.0
...
200	Porterville	200.0
201	Ventura	201.0
202	Oxnard	202.0
203	Thousand Oaks	203.0
204	Davis	204.0

205 rows × 2 columns

```
In [ ]: new_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1306 entries, 0 to 1305
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   nbh_pol                1306 non-null   object
1   med_hse_val            1306 non-null   float64
2   hse_med_age            1306 non-null   float64
3   tot_house              1306 non-null   float64
4   tot_bed                1300 non-null   float64
5   loc_pol                1306 non-null   float64
6   avg_house              1293 non-null   float64
7   avg_bed                1292 non-null   float64
8   avg_bed_per_house      1292 non-null   float64
9   ls_1_ocn               1306 non-null   float64
10  inland                 1306 non-null   float64
11  nr_bay                 1306 non-null   float64
12  nr_ocn                 1306 non-null   float64
dtypes: float64(12), object(1)
memory usage: 132.8+ KB
```

```
In [ ]: # Calculating the bedrooms per household
new_df['bed_per_hse'] = new_df['tot_bed'].div(new_df['tot_house'], axis=0)
```

```
In [ ]: new_df.isna().count()
```

```
Out[ ]: nbh_pol                1306
med_hse_val            1306
hse_med_age            1306
tot_house              1306
tot_bed                1306
loc_pol                1306
avg_house              1306
avg_bed                1306
avg_bed_per_house      1306
ls_1_ocn               1306
inland                 1306
nr_bay                 1306
nr_ocn                 1306
bed_per_hse            1306
dtype: int64
```

```
In [ ]: # Checking out the cities to add
new_df[new_df['nbh_pol']=='Brooktree'], new_df[new_df['nbh_pol']=='Fisherman's Wharf'], new_df[new_df['nbh_pol']=='Los Osos']

Out[ ]: (      nbh_pol  med_hse_val  hse_med_age  tot_house  tot_bed  loc_pol  \
      130  Brooktree      257400.0          9.0      1438.0      NaN      182.0

      avg_house  avg_bed  avg_bed_per_house  ls_1_ocn  inland  nr_bay  \
      130  532.506148  548.538144          1.0      1.0      0.0      0.0

      nr_ocn  bed_per_hse
      130      0.0      NaN ,
      nbh_pol  med_hse_val  hse_med_age  tot_house  tot_bed  loc_pol  \
      390  Fisherman's Wharf      500001.0          52.0      250.0      317.0      160.0

      avg_house  avg_bed  avg_bed_per_house  ls_1_ocn  inland  nr_bay  \
      390      501.0  535.384899          1.0      0.0      0.0      1.0

      nr_ocn  bed_per_hse
      390      0.0      1.268 ,
      nbh_pol  med_hse_val  hse_med_age  tot_house  tot_bed  loc_pol  \
      604  Los Osos      221612.5      15.375      611.75      642.5      163.0

      avg_house  avg_bed  avg_bed_per_house  ls_1_ocn  inland  nr_bay  nr_ocn  \
      604      611.75      642.5          1.0      0.0      0.0      0.0      1.0

      bed_per_hse
      604      1.050266 )
```

Ingest Data into Feature Store + Setup Feature Group

```
In [ ]: # Creating the names and the time-stamp
neighborhood_feature_group_name = "neighborhood-feature-group-" + strftime("%d-%H-%M-%S", gmtime())
encoded_feature_group_name = "encoded-feature-group-" + strftime("%d-%H-%M-%S", gmtime())
```

```
In [ ]: # Creating Feature Group
neighborhood_feature_group = FeatureGroup(name=neighborhood_feature_group_name, sagemaker_session=feature_store_session)
encoded_feature_group = FeatureGroup(name=encoded_feature_group_name, sagemaker_session=feature_store_session)
```

```
In [ ]: # Make all objects into string
convert_to_strings(new_df)
```

```
convert_to_strings(encode_df)
```

Setup Record Identifier and Event Time Features

```
In [ ]: # Creating record identifier + time features
primary_key_identifier = 'nbh_pol'
secondary_key_identifier = 'nbh_pol_new' # pol_loc
event_time_identifier = 'event_time'

# Tack on the event time to the two df
new_df[event_time_identifier] = pd.Series([time_now]*len(new_df), dtype='float64')
encode_df[event_time_identifier] = pd.Series([time_now]*len(encode_df), dtype='float64')
```

```
In [ ]: new_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1306 entries, 0 to 1305
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   nbh_pol                1306 non-null   object
1   med_hse_val           1306 non-null   float64
2   hse_med_age           1306 non-null   float64
3   tot_house             1306 non-null   float64
4   tot_bed               1300 non-null   float64
5   loc_pol               1306 non-null   float64
6   avg_house             1293 non-null   float64
7   avg_bed              1292 non-null   float64
8   avg_bed_per_house     1292 non-null   float64
9   ls_1_ocn             1306 non-null   float64
10  inland                1306 non-null   float64
11  nr_bay                1306 non-null   float64
12  nr_ocn                1306 non-null   float64
13  bed_per_hse           1300 non-null   float64
14  event_time            1306 non-null   float64
dtypes: float64(14), object(1)
memory usage: 153.2+ KB
```

```
In [ ]: encode_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 3 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   nbh_pol_new      205 non-null    object
 1   nbh_pol_encode   205 non-null    float64
 2   event_time       205 non-null    float64
dtypes: float64(2), object(1)
memory usage: 4.9+ KB
```

Load Feature Defintions

```
In [ ]: neighborhood_feature_group.load_feature_definitions(data_frame=new_df)
```

```
Out[ ]: [FeatureDefinition(feature_name='nbh_pol', feature_type=<FeatureTypeEnum.STRING: 'String'>, collection_type=None),
        FeatureDefinition(feature_name='med_hse_val', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='hse_med_age', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='tot_house', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='tot_bed', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='loc_pol', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='avg_house', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='avg_bed', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='avg_bed_per_house', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='ls_1_ocn', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='inland', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='nr_bay', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='nr_ocn', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='bed_per_hse', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='event_time', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None)]
```

```
In [ ]: encoded_feature_group.load_feature_definitions(data_frame=encode_df)
```

```
Out[ ]: [FeatureDefinition(feature_name='nbh_pol_new', feature_type=<FeatureTypeEnum.STRING: 'String'>, collection_type=None),
        FeatureDefinition(feature_name='nbh_pol_encode', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None),
        FeatureDefinition(feature_name='event_time', feature_type=<FeatureTypeEnum.FRACTIONAL: 'Fractional'>, collection_type=None)]
```

Create a Feature Group

```
In [ ]: neighborhood_feature_group.create(  
    s3_uri=f"s3://{default_s3_bucket_name}/{prefix}",  
    record_identifier_name=primary_key_identifier,  
    event_time_feature_name=event_time_identifier,  
    role_arn=role,  
    enable_online_store=True,  
)  
  
encoded_feature_group.create(  
    s3_uri=f"s3://{default_s3_bucket_name}/{prefix}",  
    record_identifier_name=secondary_key_identifier,  
    event_time_feature_name=event_time_identifier,  
    role_arn=role,  
    enable_online_store=True,  
)  
  
wait_for_feature_group_creation_complete(feature_group=neighborhood_feature_group)  
wait_for_feature_group_creation_complete(feature_group=encoded_feature_group)
```

Waiting for Feature Group Creation

Waiting for Feature Group Creation

Waiting for Feature Group Creation

FeatureGroup neighborhood-feature-group-24-20-56-29 successfully created.

Waiting for Feature Group Creation

FeatureGroup encoded-feature-group-24-20-56-29 successfully created.

Checking Results + Putting Records In

```
In [ ]: neighborhood_feature_group.describe()
```

```

Out[ ]: {'FeatureGroupArn': 'arn:aws:sagemaker:us-east-1:004608622582:feature-group/neighborhood-feature-group-24-20-56-29',
'FeatureGroupName': 'neighborhood-feature-group-24-20-56-29',
'RecordIdentifierFeatureName': 'nbh_pol',
'EventTimeFeatureName': 'event_time',
'FeatureDefinitions': [{'FeatureName': 'nbh_pol', 'FeatureType': 'String'},
{'FeatureName': 'med_hse_val', 'FeatureType': 'Fractional'},
{'FeatureName': 'hse_med_age', 'FeatureType': 'Fractional'},
{'FeatureName': 'tot_house', 'FeatureType': 'Fractional'},
{'FeatureName': 'tot_bed', 'FeatureType': 'Fractional'},
{'FeatureName': 'loc_pol', 'FeatureType': 'Fractional'},
{'FeatureName': 'avg_house', 'FeatureType': 'Fractional'},
{'FeatureName': 'avg_bed', 'FeatureType': 'Fractional'},
{'FeatureName': 'avg_bed_per_house', 'FeatureType': 'Fractional'},
{'FeatureName': 'ls_1_ocn', 'FeatureType': 'Fractional'},
{'FeatureName': 'inland', 'FeatureType': 'Fractional'},
{'FeatureName': 'nr_bay', 'FeatureType': 'Fractional'},
{'FeatureName': 'nr_ocn', 'FeatureType': 'Fractional'},
{'FeatureName': 'bed_per_hse', 'FeatureType': 'Fractional'},
{'FeatureName': 'event_time', 'FeatureType': 'Fractional'}],
'CreationTime': datetime.datetime(2024, 5, 24, 20, 56, 30, 717000, tzinfo=tzlocal()),
'OnlineStoreConfig': {'EnableOnlineStore': True},
'OfflineStoreConfig': {'S3StorageConfig': {'S3Uri': 's3://sagemaker-us-east-1-004608622582/sagemaker-featurestore-demo',
'ResolvedOutputS3Uri': 's3://sagemaker-us-east-1-004608622582/sagemaker-featurestore-demo/004608622582/sagemaker/us-east-1/offline-store/neighborhood-feature-group-24-20-56-29-1716584190/data'},
'DisableGlueTableCreation': False,
'DataCatalogConfig': {'TableName': 'neighborhood_feature_group_24_20_56_29_1716584190',
'Catalog': 'AwsDataCatalog',
'Database': 'sagemaker_featurestore'}}},
'ThroughputConfig': {'ThroughputMode': 'OnDemand'},
'RoleArn': 'arn:aws:iam::004608622582:role/LabRole',
'FeatureGroupStatus': 'Created',
'OnlineStoreTotalSizeBytes': 0,
'ResponseMetadata': {'RequestId': '5f0a3db2-6997-41ce-be82-7f490aad88f3',
'HTTPStatusCode': 200,
'HTTPHeaders': {'x-amzn-requestid': '5f0a3db2-6997-41ce-be82-7f490aad88f3',
'content-type': 'application/x-amz-json-1.1',
'content-length': '2248',
'date': 'Fri, 24 May 2024 20:56:53 GMT'}},
'RetryAttempts': 0}}

```



```
In [ ]: encoded_feature_group.describe()
```

```
Out[ ]: {'FeatureGroupArn': 'arn:aws:sagemaker:us-east-1:004608622582:feature-group/encoded-feature-group-24-20-56-29',
'FeatureGroupName': 'encoded-feature-group-24-20-56-29',
'RecordIdentifierFeatureName': 'nbh_pol_new',
'EventTimeFeatureName': 'event_time',
'FeatureDefinitions': [{'FeatureName': 'nbh_pol_new',
'FeatureType': 'String'},
{'FeatureName': 'nbh_pol_encode', 'FeatureType': 'Fractional'},
{'FeatureName': 'event_time', 'FeatureType': 'Fractional'}],
'CreationTime': datetime.datetime(2024, 5, 24, 20, 56, 32, 482000, tzinfo=tzlocal()),
'OnlineStoreConfig': {'EnableOnlineStore': True},
'OfflineStoreConfig': {'S3StorageConfig': {'S3Uri': 's3://sagemaker-us-east-1-004608622582/sagemaker-featurestore-d
emo',
'ResolvedOutputS3Uri': 's3://sagemaker-us-east-1-004608622582/sagemaker-featurestore-demo/004608622582/sagemaker/
us-east-1/offline-store/encoded-feature-group-24-20-56-29-1716584192/data'}},
'DisableGlueTableCreation': False,
'DataCatalogConfig': {'TableName': 'encoded_feature_group_24_20_56_29_1716584192',
'Catalog': 'AwsDataCatalog',
'Database': 'sagemaker_featurestore'}},
'ThroughputConfig': {'ThroughputMode': 'OnDemand'},
'RoleArn': 'arn:aws:iam::004608622582:role/LabRole',
'FeatureGroupStatus': 'Created',
'OnlineStoreTotalSizeBytes': 0,
'ResponseMetadata': {'RequestId': '62b190da-bd03-4690-800f-a1767f7729e3',
'HTTPStatusCode': 200,
'HTTPHeaders': {'x-amzn-requestid': '62b190da-bd03-4690-800f-a1767f7729e3',
'content-type': 'application/x-amz-json-1.1',
'content-length': '1583',
'date': 'Fri, 24 May 2024 20:57:01 GMT'}},
'RetryAttempts': 0}}
```

```
In [ ]: neighborhood_feature_group.ingest(data_frame=new_df, max_workers=3, wait=True)
```

```
Out[ ]: IngestionManagerPandas(feature_group_name='neighborhood-feature-group-24-20-56-29', feature_definitions={'nbh_pol': {'FeatureName': 'nbh_pol', 'FeatureType': 'String'}, 'med_hse_val': {'FeatureName': 'med_hse_val', 'FeatureType': 'Fractional'}, 'hse_med_age': {'FeatureName': 'hse_med_age', 'FeatureType': 'Fractional'}, 'tot_house': {'FeatureName': 'tot_house', 'FeatureType': 'Fractional'}, 'tot_bed': {'FeatureName': 'tot_bed', 'FeatureType': 'Fractional'}, 'loc_pol': {'FeatureName': 'loc_pol', 'FeatureType': 'Fractional'}, 'avg_house': {'FeatureName': 'avg_house', 'FeatureType': 'Fractional'}, 'avg_bed': {'FeatureName': 'avg_bed', 'FeatureType': 'Fractional'}, 'avg_bed_per_house': {'FeatureName': 'avg_bed_per_house', 'FeatureType': 'Fractional'}, 'ls_1_ocn': {'FeatureName': 'ls_1_ocn', 'FeatureType': 'Fractional'}, 'inland': {'FeatureName': 'inland', 'FeatureType': 'Fractional'}, 'nr_bay': {'FeatureName': 'nr_bay', 'FeatureType': 'Fractional'}, 'nr_ocn': {'FeatureName': 'nr_ocn', 'FeatureType': 'Fractional'}, 'bed_per_hse': {'FeatureName': 'bed_per_hse', 'FeatureType': 'Fractional'}, 'event_time': {'FeatureName': 'event_time', 'FeatureType': 'Fractional'}}, sagemaker_fs_runtime_client_config=<botocore.config.Config object at 0x7f568d93b100>, sagemaker_session=<sagemaker.session.Session object at 0x7f568d9d4730>, max_workers=3, max_processes=1, profile_name=None, _async_result=<multiprocessing.pool.MapResult object at 0x7f568d9661a0>, _processing_pool=<pool ProcessPool(ncpus=1)>, _failed_indices=[])
```

```
In [ ]: encoded_feature_group.ingest(data_frame=encode_df, max_workers=5, wait=True)
```

```
Out[ ]: IngestionManagerPandas(feature_group_name='encoded-feature-group-24-20-56-29', feature_definitions={'nbh_pol_new': {'FeatureName': 'nbh_pol_new', 'FeatureType': 'String'}, 'nbh_pol_encode': {'FeatureName': 'nbh_pol_encode', 'FeatureType': 'Fractional'}, 'event_time': {'FeatureName': 'event_time', 'FeatureType': 'Fractional'}}, sagemaker_fs_runtime_client_config=<botocore.config.Config object at 0x7f568d93b100>, sagemaker_session=<sagemaker.session.Session object at 0x7f568d9d4730>, max_workers=5, max_processes=1, profile_name=None, _async_result=<multiprocessing.pool.MapResult object at 0x7f568d9663e0>, _processing_pool=<pool ProcessPool(ncpus=1)>, _failed_indices=[])
```

```
In [ ]: ### Grabbing the Record from the online store
record_identifier_value = 'Brooktree'

featurestore_runtime.get_record(
    FeatureGroupName=neighborhood_feature_group_name,
    RecordIdentifierValueAsString=record_identifier_value,
)
```

```
Out[ ]: {'ResponseMetadata': {'RequestId': '524106e0-3df5-48df-bfe4-0b5ca3f12fc3',
    'HTTPStatusCode': 200,
    'HTTPHeaders': {'x-amzn-requestid': '524106e0-3df5-48df-bfe4-0b5ca3f12fc3',
    'content-type': 'application/json',
    'content-length': '1054',
    'date': 'Fri, 24 May 2024 21:01:11 GMT'},
    'RetryAttempts': 0},
    'Record': [{'FeatureName': 'nbh_pol', 'ValueAsString': 'Brooktree'},
    {'FeatureName': 'med_hse_val', 'ValueAsString': '257400.0'},
    {'FeatureName': 'hse_med_age', 'ValueAsString': '9.0'},
    {'FeatureName': 'tot_house', 'ValueAsString': '1438.0'},
    {'FeatureName': 'loc_pol', 'ValueAsString': '182.0'},
    {'FeatureName': 'avg_house', 'ValueAsString': '532.5061475409836'},
    {'FeatureName': 'avg_bed', 'ValueAsString': '548.5381443298969'},
    {'FeatureName': 'avg_bed_per_house', 'ValueAsString': '1.0'},
    {'FeatureName': 'ls_1_ocn', 'ValueAsString': '1.0'},
    {'FeatureName': 'inland', 'ValueAsString': '0.0'},
    {'FeatureName': 'nr_bay', 'ValueAsString': '0.0'},
    {'FeatureName': 'nr_ocn', 'ValueAsString': '0.0'},
    {'FeatureName': 'event_time', 'ValueAsString': '1716584186.0'}]}
```

```
In [ ]: record_identifier_value = "Fisherman's Wharf"

featurestore_runtime.get_record(
    FeatureGroupName=neighborhood_feature_group_name,
    RecordIdentifierValueAsString=record_identifier_value,
)
```

```
Out[ ]: {'ResponseMetadata': {'RequestId': '00c5eaad-81ae-41c5-a99d-c4084c43bde',
    'HTTPStatusCode': 200,
    'HTTPHeaders': {'x-amzn-requestid': '00c5eaad-81ae-41c5-a99d-c4084c43bde',
    'content-type': 'application/json',
    'content-length': '1204',
    'date': 'Fri, 24 May 2024 21:01:13 GMT'},
    'RetryAttempts': 0},
    'Record': [{'FeatureName': 'nbh_pol', 'ValueAsString': "Fisherman's Wharf"},
    {'FeatureName': 'med_hse_val', 'ValueAsString': '500001.0'},
    {'FeatureName': 'hse_med_age', 'ValueAsString': '52.0'},
    {'FeatureName': 'tot_house', 'ValueAsString': '250.0'},
    {'FeatureName': 'tot_bed', 'ValueAsString': '317.0'},
    {'FeatureName': 'loc_pol', 'ValueAsString': '160.0'},
    {'FeatureName': 'avg_house', 'ValueAsString': '501.0'},
    {'FeatureName': 'avg_bed', 'ValueAsString': '535.3848987108655'},
    {'FeatureName': 'avg_bed_per_house', 'ValueAsString': '1.0'},
    {'FeatureName': 'ls_1_ocn', 'ValueAsString': '0.0'},
    {'FeatureName': 'inland', 'ValueAsString': '0.0'},
    {'FeatureName': 'nr_bay', 'ValueAsString': '1.0'},
    {'FeatureName': 'nr_ocn', 'ValueAsString': '0.0'},
    {'FeatureName': 'bed_per_hse', 'ValueAsString': '1.268'},
    {'FeatureName': 'event_time', 'ValueAsString': '1716584186.0'}]}
```

```
In [ ]: record_identifier_value = 'Los Osos'

featurestore_runtime.get_record(
    FeatureGroupName=neighborhood_feature_group_name,
    RecordIdentifierValueAsString=record_identifier_value,
)
```

```
Out[ ]: {'ResponseMetadata': {'RequestId': '010e5a93-7e39-4eca-95e4-337f31c914aa',
    'HTTPStatusCode': 200,
    'HTTPHeaders': {'x-amzn-requestid': '010e5a93-7e39-4eca-95e4-337f31c914aa',
    'content-type': 'application/json',
    'content-length': '1200',
    'date': 'Fri, 24 May 2024 21:01:15 GMT'},
    'RetryAttempts': 0},
'Record': [{'FeatureName': 'nbh_pol', 'ValueAsString': 'Los Osos'},
{'FeatureName': 'med_hse_val', 'ValueAsString': '221612.5'},
{'FeatureName': 'hse_med_age', 'ValueAsString': '15.375'},
{'FeatureName': 'tot_house', 'ValueAsString': '611.75'},
{'FeatureName': 'tot_bed', 'ValueAsString': '642.5'},
{'FeatureName': 'loc_pol', 'ValueAsString': '163.0'},
{'FeatureName': 'avg_house', 'ValueAsString': '611.75'},
{'FeatureName': 'avg_bed', 'ValueAsString': '642.5'},
{'FeatureName': 'avg_bed_per_house', 'ValueAsString': '1.0'},
{'FeatureName': 'ls_1_ocn', 'ValueAsString': '0.0'},
{'FeatureName': 'inland', 'ValueAsString': '0.0'},
{'FeatureName': 'nr_bay', 'ValueAsString': '0.0'},
{'FeatureName': 'nr_ocn', 'ValueAsString': '1.0'},
{'FeatureName': 'bed_per_hse', 'ValueAsString': '1.0502656313853698'},
{'FeatureName': 'event_time', 'ValueAsString': '1716584186.0'}]}
```

```
In [ ]: a
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[45], line 1
----> 1 a

NameError: name 'a' is not defined
```

```
In [ ]: %html

<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Ker

<script>
try {
    els = document.getElementsByClassName("sm-command-button");
    els[0].click();
}
```

```
catch(err) {  
    // NoOp  
}  
</script>
```